

Significant updates since initial Design Doc submission:

Performance Testing:

Our OS boots with optimized compiling with as little as 472K RAM with 1 CPU, and up to 10 CPUs at 512K RAM (though barely anything runs at these configurations). We tested primarily with a 128M swap disk (though it can be arbitrarily large or small, within reason), but did check that storage wasn't being leaked using the kernel menu command `cm`, which prints the core map and whether swap is completely zeroed.

Our TLB eviction policy is random (but OS-level random, not TLB's random that excludes the first 8 entries). We deemed the extra cost of traversing the TLB's entries and storing/performing calculations and heuristics not worth the marginal gains one would get from doing so. There are much more obvious performance optimizations for the TLB that would offer definitive improvements rather than the workload-based tradeoff that a different eviction algorithm would entail. Supporting ASIDs would partially eliminate the cost of cold misses when resuming execution of a thread, and a fast-path TLB handler would cut the number of instructions required to refill the TLB by several orders of magnitude (because it would circumvent the entire trap handling procedure). Finally, our page eviction algorithm minimizes the risk that unwise TLB evictions could lead to highly demanded pages being swapped out, so its performance is not as important as it otherwise might be.

Our page eviction/swap policy is a version of the clock algorithm (itself an approximation of the least recently used policy) that iterates over the core map and loops from the final to first element. The clock hand starts at a persistently updated position and keeps looping until a suitable page to swap out is found (or after three revolutions it stops, meaning only kernel or busy pages are left in the core map). On the first loop it prefers user pages that are not in the TLB and were not recently in the TLB, marking those recent pages as not-recent when skipping over them. On the second loop it now has access to those previously "recent" pages. On the third loop it begins to allow selection of pages in a TLB, which is unideal because of the overhead of a TLB shutdown and the likelihood that a page in the TLB is being actively used. Even though our TLB eviction policy is random, if a shot-down page is part of a process's working set, it should soon be refilled into the TLB; the first revolution of the clock gives this time to happen without incurring cost from swapping to disk. The revolution of the clock also helps to prevent situations where a page that has just been swapped in is immediately swapped out again. We didn't focus substantially on combating antagonistic memory access patterns, but the way we

implemented things like `as_copy` (where each page that has been copied is now swappable and not recent) means that forking processes don't get to take over RAM and force other processes to thrash.

Our write-back daemon starts at the current clock position whenever it wakes up so that it can preemptively write whatever will be needed the soonest. It has similar criteria to our eviction algorithm, but it never writes back pages in the TLB; performing a TLB shutdown to remove the writeable bit for something entirely preemptive would be a bit wasteful. The daemon is intended to be active most when free pages are running out and when the number of pages in swap is a low multiple of the number of pages in physical memory. The other sides of this distribution (plenty of free memory or thrashing) are not conducive to preemptive writes to swap.

As we understand it from dholland, this section's discussion of performance tuning is mostly hypothetical. So, in addition to stats that we do track and use (for the write-back daemon) like total number of pages, number of free pages, number of dirty pages, and number of pages in swap, if we were to focus on performance optimization we'd add tracking for permissions faults, TLB misses, swap accesses, TLB shutdowns, lost TLB entries from context switches, and the time each of these events takes. When combined with the results from different test programs representing different workloads (intensive sequential or parallel calculations on a large data set, calculations using a small set of pages with high frequency and a large set of pages with low frequency, repeated forking, repeated execing, backwards array access, random array access, strided array access, etc.) would provide a much better holistic picture of our algorithms' performance.

In terms of our specific algorithms, random TLB eviction can't really be tuned, and the clock algorithm is difficult to adjust with any sort of granularity because it's largely based on binary policies. However, we kept these algorithms isolated in their own functions so that alternatives could easily be tested, with the active policy determined by a config file. The write-back daemon is easier to "tune" because it has parameters like sleep time, number of dirty writes per sleep, etc. that can be adjusted to fit what we think its behavior should look like at different levels of memory pressure, but in actuality when testing we found that these parameters had relatively little effect on actual performance, and if they had any effect, it was usually a negative one. We settled on a configuration that seemed relatively inoffensive.

Final structs:

addrspace.h

```
struct addrspace {
    struct page_table_directory* ptd;
    struct spinlock addr_splk;
    struct wchan *addr_wchan;
    vaddr_t heap_bottom;
    vaddr_t heap_top;
};
```

machine/vm.h

```
union page_table_entry {
    struct {
        unsigned int addr : 20, : 7;    // address in memory or swap
        unsigned int x : 1;              // executable (unused)
        unsigned int r : 1;              // readable (unused)
        unsigned int w : 1;              // writeable (unused)
        unsigned int p : 1;              // present
        unsigned int b : 1;              // busy
    };
    uint32_t all;    // for zeroing the PTE in one instruction
};
```

```
struct page_table {
    union page_table_entry ptes[NUM_PTES];
};

struct page_table_directory {
    struct page_table* pts[NUM_PTES];
};
```

```
struct tlbshootdown {
    uint32_t oldentryhi;
    struct addrspace *as;
};
```

vm.h

```
union metadata {
    struct {
        unsigned int swap : 20, : 5;    // address in swap
        unsigned int recent : 1;         // recently evicted from TLB
        unsigned int tlb : 1;           // currently in TLB
        unsigned int dirty : 1;         // dirty page
    };
};
```

```
        unsigned int contig : 1;        // end of kernel allocation
        unsigned int kernel : 1;        // belongs to kernel
        unsigned int s_pres : 1;        // present in swap
        unsigned int busy : 1;          // busy
    };
    uint32_t all;        // for zeroing metadata in one instruction
};
```

```
struct core_map_entry {
    vaddr_t va;          // virtual address of the page
    struct addrspace *as; // address space for the virtual address
    uint32_t reserved;    // unused
    union metadata md;    // 4 bytes of metadata
};
```

Computer Science 161: Operating Systems

Assignment 3 Design Document¹

David Li and Garrett Tanzer
March 10, 2017

1. Introduction

Assignment 3 can be broken up into three major data structures (page tables hierarchies, the core map, and swap) and the operations performed to maintain them (populating/evicting entries from the TLB, swapping pages from disk to RAM as needed, and asynchronously writing dirty pages to disk with a daemon).

We will implement 2-level page tables, with each level indexed by 10 bits of the virtual address; 12 physical bits are then added to the end to get a 32-bit physical address. The core map will be an array of entries for each page of physical RAM with its own bookkeeping information. Swap, twice the size of physical RAM (rounded up to the nearest page) will be tracked using a bitmap. Each address space, the core map, and the swap bitmap will be protected by a coarse-grained spinlock. Each entry of the core map and PTE is also protected by a busy bit, which can be set while coarse synchronization is held for each, in order to prevent accesses during the swap process.

TLB eviction will be largely random. Page swapping will be a variant of the clock algorithm, dependent in part on presence in the TLB. The write-back daemon will loop over the core map sequentially and save everything in RAM to disk when the pages are not the latest version. The main challenge with these operations performing in parallel is synchronization; our algorithms and the primitives we use to ensure they are performed atomically are explained in much greater detail below.

2. Overview

¹ Based on this Piazza post (<https://piazza.com/class/iy1oku7camk18k?cid=159>), our current design document does not have a section explicitly outlining how performance optimization will happen because it is supposedly only needed in the final submission. In addition to simply running testbin programs to fine-tune parameters of TLB eviction, page swapping, and the write-back Daemon, we can add stat-tracking to see how many page faults, TLB misses, swaps, etc. each strategy has, and we could even time how long it takes to resolve each of them. This collection of metrics would more facilitate much more meaningful analysis than the time to run a certain program.

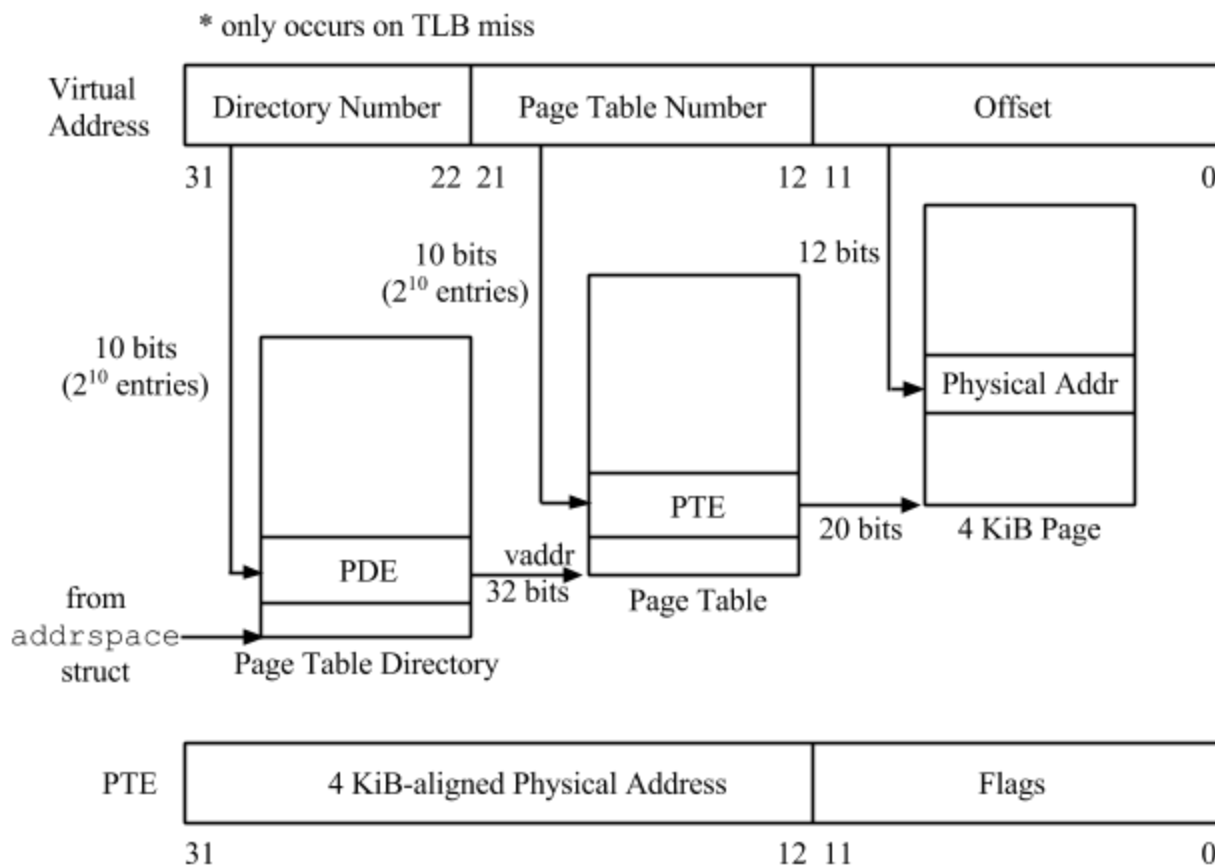
2.1 Page Tables

The data structures for virtual to physical address translation begin in the `addrspace` struct:

```
struct addrspace {
    struct page_table_directory* ptd;
    struct spinlock addr_splk;
    struct wchan *addr_wchan;           // for waiting during swap I/O
    vaddr_t heap_bottom;
    vaddr_t heap_top;
};
```

`ptd` provides an entry point into the hierarchical page tables, while `addr_splk` protects modifications to this address space. `heap_top` is manipulated by `sbrk()` and is used to check for invalid memory accesses.

Here is a diagram showing how two-level page tables are indexed to reach a full 32-bit physical address:



The `ptd` field of `addrspace` is a wrapper for an array of `page_table` pointers (physical addresses of subsequent page tables or NULL if not yet created). Each `page_table` has an

array of 4-byte entries, which store a 20-bit physical frame and 12 bits of flags. Both page table levels are designed to take exactly one page of memory each so that they can fill exactly one use of `alloc_kpages()`.

```
struct page_table_directory{
    struct page_table* pts[1024];
};
```

```
struct page_table {
    uint32_t ptes[1024];
};
```

Each of the PTEs in `page_table` is organized as follows:

PADDR/SWAP								PTE_X	PTE_R	PTE_W	PTE_B	PTE_P
31-12	11	10	9	8	7	6	5	4	3	2	1	0

PTE_P: 1 if present in memory, 0 if in swap

PADDR: If PTE_P, the physical address is stored in the top 20 bits.

SWAP: If not PTE_P, the index of the page in swap is stored in the top 20 bits.

PTE_B: 1 if busy, 0 if not

PTE_W: 1 if writable, 0 if not

PTE_R: 1 if readable, 0 if not

PTE_X: 1 if executable, 0 if not

Rather than having to manually traverse the page table with

`curproc->p_addrspace->ptd->pts[index1]->ptes[index2]`, we will define a function `as_pte(vaddr_t vaddr)` that returns the PTE mapped by the virtual address parameter in the current address space (or a negative error value if it has not yet been mapped).

2.2 Core Map

The core map is a global array with bookkeeping data on all page frames in RAM. It will have `ram_getsize() / PAGE_SIZE` entries and must be initialized early in the bootstrap process using `ram_stealmem()`, since the core map is required for `kmalloc()` to work. We will need to mark the RAM used by the core map in the core map itself. The entire array will be protected by a single spinlock. Each entry will be protected by a busy bit.

```

struct core_map_entry {
    vaddr_t va;
    struct addrspace *as;
    uint32_t metadata;
};

```

`va` and `as` are used to locate the page table entry pointing to the physical address, as well as to acquire the spinlock for the address space.

Each metadata in the core map is structured as follows:

Our final code will probably use bit fields to represent this instead for simpler flag access.

SWAP	RECENT	TLB	DIRTY	CONTIG	KERNEL	BUSY	S_PRES
12-31	11-8	7	6	5	4	3	2	1	0

S_PRES: 1 if copy already exists in swap, 0 if not
BUSY: 1 if this page is involved in swapping right now, 0 if not
KERNEL: 1 if this page belongs to the kernel, 0 if to the user level
CONTIG: 1 if this page is the end of a contiguous kernel allocation, 0 if not
DIRTY: 1 if this page has been written to since the last flush to swap, 0 if not
TLB: 1 if this page's translation is currently in a TLB, 0 if not
RECENT: 1 if recently evicted from a TLB, 0 if not
SWAP: Index of the corresponding page on the disk: Assuming a max swap size of 1 GiB (2 * 512 MiB max RAM in kseg0)

2.3 Swap

The swap space is a region of the hard drive that stores pages either already evicted from RAM or preemptively copied from RAM, in anticipation of the page needing to be evicted from RAM. We will keep track of occupied/free sectors on the swap device using a `struct bitmap` of length `2 * ram_getsize() / PAGE_SIZE` bits, rounded up to the nearest page. So, with physical RAM limited to 512 MiB, this means our swap device can be between 128 and 1024 MiB inclusive. This space was made larger than RAM to ensure that even if all the pages are preemptively sent to swap, there will still be space for others to actually be swapped in/out.

The bitmap will be protected by a single spinlock. Whenever a new swap index is needed, the first available 0 in the bitmap is switched to 1 under the spinlock's protection.

Read/write to swap using the vnode generated by `vfs_swapon("lhd0raw")`.

3. Topics

3.1 TLB Misses

The TLB miss handler needs to populate the TLB with the translation for a requested virtual page. This apparently simple task becomes exceedingly complicated when you consider swap and synchronization. Not only must the handler respond to requests for invalid addresses, but for valid requests it must make sure the requested page is in memory (and if not, swap it in, which may require swapping out another page if memory is full), writing that translation to the TLB, and updating bookkeeping data.

Our specific pseudocode implementation is below, but in general terms our strategy for avoiding deadlock and race conditions is to only acquire spinlocks in the order address space > core map > swap and to release those spinlocks as soon as possible (using busy bits in PTEs and the core map for loosely enforced polling synchronization afterwards).

When `VM_FAULT_READ` or `VM_FAULT_WRITE` are triggered:

- If the faulting address is not in user space:
 - Kill the thread with a segfault
- Acquire the address space's spinlock
- If the faulting address is unmapped but a valid address (i.e. in stack/heap bounds):
 - Call `alloc_upages` to allocate a new page
- If the faulting address is mapped, but it's read only and you got `VM_FAULT_WRITE`:
(You know that once you populate the TLB with the new entry it'll just fail when the instruction runs again, so why not catch it now and save the time?)
 - Release the address space spinlock
 - Kill the thread with a segfault
- While `PTE_B` is 1:
 - Sleep on the original address space's wchan with its spinlock
- If the `PTE_P` is not set
 - Set `PTE_B` to 1
 - Release the address space lock
 - Acquire the core map spinlock
 - Iterate through core map, looking for free entries
 - If free entry found:
 - Save the index of the core map entry
 - Else:
 - // swap in/out will probably be helper functions in the final code, but it's easier to follow synchronization if it's included here
 - // there will probably be an optional parameter to leave pages BUSY

- Iterate through the core map starting from a global index variable `clock`, which loops back to the beginning of the array when it reaches the end
 - If `TLB` or `BUSY`, continue (as in skip to the next iteration of the loop)
 - If `RECENT`, set `RECENT` to 0 and continue
 - On the second revolution through the clock, start accepting entries with a set `TLB` flag
 - If you reach this point in the loop without hitting “continue,” choose the current index to swap out to disk
- Set the `BUSY` bit on your chosen core map entry
- Release the core map lock
- Acquire the new address space spinlock (for the core map entry that you want to evict)
- Set `PTE_B` of the PTE referencing the core map entry
- If the `TLB` bit is set:
 - Broadcast a TLB shutdown
 - Go to sleep on the new address space’s `wchan` and its spinlock (TLB shutdown will wake all on its own `addrspace`’s `wchan` when it’s done)
- Release the new address space spinlock
- If `S_PRE` is set:
 - Set `PTE_P` to 0 and set `SWAP` index in the PTE to `SWAP` from the core map entry
 - If the core map entry is `DIRTY`:
 - Write the contents of the physical page to disk
- Else (the cache needs to be written to RAM):
 - Acquire the swap spinlock
 - Find a free spot in the bitmap
 - Mark it taken
 - Release the swap spinlock
 - Set `PTE_P` to 0 and set `SWAP` index in the PTE to the free spot
 - Write the contents of the physical page to disk
- Set `PTE_B` of the PTE to 0
- Acquire the core map lock
- Wake all on the new address space’s `wchan`
- Acquire original address space spinlock
- Set the core map entry’s `SWAP` to the value from the original (faulting) PTE
- Set the core map entry’s virtual address to the faulting virtual address
- Set the core map entry’s address space to the faulting address space

- Put the physical address associated with the core map entry into the PTE
- Release the original address space spinlock
- Read data from disk at `SWAP` into the physical page of memory
 - We don't need to worry about the PTE being modified in this period because OS/161 is single-threaded, so any other attempted access to this specific PTE will first need to come through the core map entry, which is tagged as `BUSY`
- Acquire the original address space spinlock
- Set `PTE_B` on the original address space spinlock to 0
- Set `BUSY` to 0 on the core map entry (because of the branches above it is possible that this was already 0, but it is impossible that another thread set this `BUSY`)
- Wake all on the old address space's `wchan`
- Save the index of the core map entry
- Holding a spinlock means interrupts are already off, so we don't need to worry about TLB state getting messed up
- Set `TLB` on the core map entry to 1
- Use OS-level random function to determine which TLB entry to evict (because `tlb_random()` only selects 56 of 64 entries)
- Read this randomly chosen entry from the TLB with `tlb_read()`
- While this randomly chosen entry's core map is `BUSY`:
 - Choose a different random one—dealing with someone swapping right now would be annoying
- Set `TLB` to 0 and `RECENT` to 1 for the randomly chosen entry's core map
- Write the replacement entry to the TLB with `tlb_write()`, with global and valid bits set (not writeable)
- Release core map spinlock
- Vector control back to the user level, with the instruction pointer at the faulting instruction

3.2 Page Faults

The page fault handler should respond to accesses with improper permissions. Sometimes these will actually be illegal accesses according to the PTE, in which case we should kill the offending thread, but most of the time these will be intentional faults in order to set the `DIRTY` bit for the swap write-back daemon. In this case, we mark the relevant bookkeeping information and update the TLB to allow writes so that the original instruction can go through.

When `VM_FAULT_READONLY` is triggered: (should only happen for valid TLB entries, but not writeable)

- Acquire the address space spinlock
- If the `PTE_W` isn't set in the PTE:
 - Release the address space spinlock
 - Kill the thread with a segfault
- Acquire the core map spinlock
- If the physical page is `BUSY`:
 - Release the core map spinlock
 - While the physical page is `BUSY`:
 - Go to sleep on the address space wait channel with its spinlock
 - Acquire the core map spinlock
- Mark the associated physical page's `DIRTY` bit
- Again, the core map spinlock turns interrupts off for us
- Update the TLB entry for the faulting address with a writeable bit using `tlb_write()`
- Release the core map spinlock
- Release the address space spinlock
- Vector control back to user space, with the faulting instruction repeated

3.3 MAT Daemon

The MAT (Memory Access Traversal (and Write-Back)) Daemon is a background process spawned during boot to proactively write dirty pages from RAM to the swap space. It will also update pages on the swap if they have already been placed onto swap. This kernel-mode process should be given a low priority in the scheduler since it is not crucial to the functioning of the operating system. MAT Daemon will `thread_yield()` after a certain amount of time, and because it uses spinlocks, it has to give up control voluntarily; the amount of time it takes to do so will be tuned based on performance once implemented. MAT Daemon will not run if no swap space is left. Below is the algorithm for the function:

- Acquire the core map spinlock
- Iterate through core map continually and backwards, from the greatest entry of the core map to the least entry of the core map, restarting after the lowest entry is hit. This backwards accessing is used to ease thrashing between the MAT Daemon and our clock swapping algorithm.
 - If the entry is `KERNEL` or `TLB`, continue
 - If the core map entry is `DIRTY`:
 - Set `BUSY` to 1
 - Release the core map spinlock
 - If not `S_PRE`: (need to find a swap index)
 - Acquire the swap spinlock

- Find first free spot in the bitmap
 - Mark it taken
 - Release the swap spinlock
 - Set SWAP index in the core map entry to the free spot
- Set DIRTY to 0
- Write the contents of the physical page to disk at SWAP index
- Set BUSY to 0
- Acquire the core map spinlock
- At some point decide to yield to the scheduler
 - Release the core map spinlock
 - thread_yield()

3.4 TLB Shootdown

The TLB shutdown handler will use `tlb_probe()` to find the requested entry then invalidate it. Once it finishes, it must wake all on its address space's wchan (to signal to the thread that issued the shutdown that it completed). Use the constants/masks defined in `tlb.h` to construct and make sense of TLB entries.

The struct `tlbshutdown` in `vm.h` will be defined as follows:

```
struct tlbshutdown {
    uint32_t oldentryhi;    // tlb_probe() doesn't need entrylo
};
```

3.5 Miscellaneous Functions

```
void* sbrk(intptr_t amount);
```

Extends the top of the heap for the address space (i.e. `heap_top`) by the `amount`, updating the page tables accordingly. We should create a constant to indicate the greatest extent of heap expansion.

- Get address space and take the spinlock
- Check page-alignment of `amount`

If `amount` is positive:

- Test for ENOMEM condition
- Increase `heap_top`

- Release address space spinlock

If amount is negative:

- Test for `EINVAL` condition using `heap_bottom`
- Get core map's spinlock
- Decrease `heap_top`
- Zero out pages in core map corresponding to virtual addresses between old and new `heap_top`, marking associated swap pages as free in the process
- Mark associated pages only in swap as available with the swap spinlock
- Remove entries marked `TLB` from the `TLB` (this is in the same address space, so since we're multithreaded, that means it's our `TLB`, which makes it much easier)
- Release core map spinlock

Return pointer to original `heap_top`

Errors:

- `EINVAL` for an amount that is either not page-aligned or would make `heap_top` below `heap_bottom`
- `ENOMEM` for insufficient virtual memory

```
struct addrspace* as_create(void);
```

Create a new empty address space (create minimum page table entries, initialize spinlock and `wchan`, `NULL` everything else). Return `NULL` if out of memory.

```
int as_copy(struct addrspace *src, struct addrspace **ret);
```

Do a deep copy of the `src` address space into a new address space saved to `ret`. During the process, the `src` spinlock must be held.

- First, `as_create()` (this might fail, in which case return the error)
- Acquire the `src` spinlock
- Iterate through the entire address space and for each mapped entry:
 - Get one new user page with `alloc_upages()` (this acquires the new address space's lock, but because it's not yet running on a thread, there is no worry about deadlock from holding the spinlocks for two address spaces)
 - If the PTE in `src` is in memory
 - `memcpy` from it into the new page
 - If the PTE is in swap

- Read from swap directly into the new memory
- Release the address space spinlock
- Return success.

```
void as_activate();
```

Flush the TLB. Update the TLB flags in core map entries while doing so (with spinlock protection).

```
void as_deactivate();
```

Do nothing.

```
void as_destroy(struct addrspace *as);
```

Iterate through the address space and free all remaining referenced pages (using helper functions). Clean up the address space's spinlock, wchan, etc. Remember to kfree the addresses of the pagetables themselves; not just clear what they refer to.

```
int as_define_region(struct addrspace *as, vaddr_t vaddr, size_t sz, int readable, int writeable, int executable);
```

Map the designated virtual memory region to physical memory with the specified permissions in user space.

- Align `vaddr` and `sz` based on `PAGE_FRAME` and `PAGE_SIZE`
- Use `alloc_upages()` (described below) to allocate an aligned `sz / PAGE_SIZE` pages of user memory in `as` starting at `vaddr`
- Iterate through the allocated pages in virtual memory and set their permissions to the specified ones

```
int as_prepare_load(struct addrspace *as);
```

Do nothing.

```
int as_complete_load(struct addrspace *as);
```

Allocate the bottom of the heap, and set up the top in `addrspace` so that `sbrk()` will work later.

```
int as_define_stack(struct addrspace *as, vaddr_t *stackptr);
```

Allocate space and populate page tables from the top of user address space down to the given stack pointer.

```
vaddr_t alloc_kpages(unsigned pages);
```

Helper function that allocates `pages` contiguous pages of kernel memory in the core map and returns a pointer to the first one (use `PADDR_TO_KVADDR`). Make sure to use the `CONTIG` bit to mark where the allocation starts and stops and mark pages as `KERNEL`.

- Get core map lock
- Keep track of:
 - Max number of contiguous free pages so far
 - Index of that page chain
 - Max number of contiguous non-KERNEL non-TLB non-BUSY pages so far
 - Index of that page chain
 - Max number of contiguous non-KERNEL non-BUSY pages so far
 - Index of that page chain
- Iterate over core map, counting the above stats
 - If at any point you reach the required number of contiguous free pages, break
- If the number of contiguous free pages is sufficient, use that index
- Else if the number of contiguous non-KERNEL, non-TLB, non-BUSY pages is sufficient:
 - Swap the relevant pages out of memory
 - Remember the starting index
- Else if the number of contiguous non-KERNEL non-BUSY pages is sufficient:
 - Swap the relevant pages out of memory
 - In the process, TLB shutdown the relevant pages and sleep on `wchan` until done
 - Remember the starting index

- Else return ENOMEM
- Iterate over the chain of free pages and mark them as belonging to the kernel
- Convert PADDR to VADDR with the macros and return the VADDR value

```
void free_kpages(vaddr_t addr);
```

Helper function that frees the kernel allocation starting at `addr` (in `kseg0`). Remember to use the `CONTIG` bits to free the entire allocation.

```
int alloc_upages(struct addrspace *as, vaddr_t vaddr, unsigned
pages);
```

Helper function that allocates `pages` pages of user memory in the core map and populates `as`'s page tables with them starting at `vaddr`. Does not mark pages `KERNEL`.

You must not hold any locks for potentially active address spaces or the core map lock when calling `alloc_upages()`.

- Acquire address space lock
- Acquire core map lock
- Set `counter = ceil(sz/PAGESIZE)`
- Loop through every entry in core map
 - If entry is empty (has no `vaddress`)
 - Map it to the appropriate virtual page, adding more page table entries/page tables if necessary; set the permissions appropriately
 - Set the `vaddr` and `as` fields in `coremap`
 - Decrease `counter` for every successful addition; break when `counter == 0`
- Release core map lock
- Release address space lock
- While `counter > 0`:
 - Swap out a page (Refer to the `PTE_P` not set code in the paging section)
 - Acquire original address space lock
 - Acquire core map lock
 - Attempt to mark that page as ours and put it into our page table, as in the body of the for loop
 - Decrement `counter`

- Release core map lock
- Release address space lock
- Because of the way the locking works, it might be possible for another thread to swoop in and claim the page we just swapped out from, but that's okay because `counter` only decreases upon success
- Return success, unless we ran out of both memory and swap space at some point during the function

```
void free_upages(struct addrspace *as, vaddr_t vaddr, unsigned
pages);
```

Helper function that frees (from the core map) `pages` contiguous pages of user memory in the virtual address space `as`, starting at address `vaddr`. This function will also need to shoot down the associated TLB entries and deal with removing pages that were swapped out.

4. Plan of Action

Monday: core map
Tuesday: core map bootstrap
Wednesday: page tables
Thursday: `as_create()`, `as_destroy()`, `as_copy()`,
`as_prepare_load()`, `as_define_region()`, `as_complete_load()`,
`as_define_stack()`
Friday: Paging
Saturday: `as_activate()`, `as_deactivate()`, `tlbshootdown()`
Sunday: TLB Misses (No swap yet)
Monday: Swapping
Tuesday: MAT Daemon
Wednesday: `sbrk()`
Thursday: test
Friday: test

5. Code-Reading Questions

5.1 Design Questions

1. TLB miss, page fault

- page not present in TLB upon attempted read or write (can be in RAM, in swap, or nonexistent)

TLB miss, no page fault

- not possible, unless you consider use of kseg0 without the TLB a “miss”

TLB hit, page fault

- page present in TLB, wrong permissions (e.g. write to read-only address)

TLB hit, no page fault

- normal kuseg memory access (or kseg2 I suppose, but we’re not really dealing with that)

2. It reads the faulting instruction, puts the address in the TLB, reruns the instruction, puts the instruction back in the TLB, etc. Infinite loop.

3. Requested address is invalid (Negative or not aligned to word), TLB entry does not exist, Process has no permission for the memory region it’s trying to access, It’s write only

5.2 Malloc Questions

1. a. malloc only calls sbrk once throughout the loop

b. 0x21000

2. The last allocation does not need sbrk. The allocation pointed to by x is on the same page or same page range as the first allocation.

3. A user program might want to define a function of the same name; identifiers with underscore prefixes are reserved for the implementation of the C dev environment to prevent collisions with names used by applications.

4. Allocations are rounded up to an integral number of blocks. Non-fundamental data types are aligned to their greatest member. The sizes of the greatest members is a multiple of where each block starts.

5.3 Paging Algorithms

We can/will separate each eviction strategy into a separate function, then define a series of macros with the same text that each replace to a different eviction strategy. We would then

comment out all but one and compile. Alternatively, instead of macros we could make an array of function pointers that we index into using a `#defined` index that can be easily changed. If we wanted to use different structs/stats for each algorithm, we could define them all with `#if`, `#endif`, and such but that would get very convoluted.

The above strategy is in our opinion the reasonable one, but if we just *need* instantaneous switching (these answers are more speculative):

- Maybe there's a way to set a flag in `sys161.conf` and have the OS choose its swapping algorithm based on it at startup (so we wouldn't need to recompile to switch page swapping algorithms), but that would probably require deeper knowledge of the inner workings of `sys161`.
- Theoretically we could also make a kernel program that can change a synchronized variable that serves the same purpose as a `#defined` index (and that way we could switch the page swapping algorithm once the OS is running), but then they would have to all use the same struct/stats. Many algorithms might also depend on data structures that are cultivated from boot, so they wouldn't be compatible.